

Summer Internship Programme

Henry Harvin Education India LLP
Sector-2, Noida, U.P.-201306



Project Title – **Loan Prediction**

Mentor Name: Ms. Pooja Gupta (Senior Consultant)

Name: Saksham Sood

Course: Summer Internship Programme (SIP) Python

Batch: Jun-Jul 2019

Job: Business Analyst Associate (Intern)

Institution: Lovely Professional University, Jalandhar, Punjab.

DECLARATION

I hereby declare that the project report entitled “**Loan Prediction**” submitted by me to **HENRY HARVIN EDUCATION INDIA** is a record of bonafide project work carried out by me under the guidance of MS. POOJA GUPTA. This project is an original report with references taken from websites and help from mentors and teachers.

DATE: 28 Jul 2019

Saksham Sood

SIP – Python

Acknowledgements

In the accomplishment of this project successfully, many people have best owned upon me their blessings and the heart pledged support, this time I am utilizing to thank all the people who have been concerned with this project. I would like to thank my teachers MR. DHIRAJ UPADHYAYA and MR. ANIL JADON whose valuable guidance has been the ones that helped me patch this project and make it full proof success.

Their suggestions and instructions have served as the major contributor towards the completion of the project. I would like to thank my mentor MS. POOJA GUPTA for giving me this golden opportunity.

Then I would like to thank my parents and friends who have helped me with their valuable suggestions and guidance has been very helpful in various phases of the completion of the project. Last but not the least I would like to thank my batchmates who have helped me a lot.

SAKSHAM SOOD

SIP-Python

Introduction

Predicting the outcome of a loan is a recurrent, crucial and difficult issue in insurance and banking. The objective of our project is to predict whether a loan will default or not based on objective financial data only.

This is a Python-based Project. This project was created via Spyder 3.3.5. IDE (Integrated Development Environment) using Python 3.7.3 and Ipython Console 7.4.0. The final outcome of this project is saved in a Jupyter Notebook v7.8.0. The libraries of python used in this project are:

1. NumPy
2. Pandas
3. Matplotlib
4. Seaborn
5. Stats models
6. Sci-kit Learn

This project is based on a data set provided by the teachers via GITHUB.

Since the objective is to predict the outcome from the information gathered at the signature of the loan, we cannot use the data concerning the history of payments or the current situation of a loan.

Excluding features for which the information is incomplete, or uninformative, we get a total of 19 features, that cover personal information (credit grade, income, housing status, ...) and credit information (amount, interest rates, ...). Accuracy is not well-suited for our problem. The unbalance of the classes would lead an algorithm to never predict a default. F1-score allows us to quantify a good prediction on both precision and recall.

Loan Prediction

Problem

- A Company wants to automate the loan eligibility process (real time) based on customer detail provided while filling online application form. These details are Gender, Marital Status, Education, Number of Dependents, Income, Loan Amount, Credit History and others. To automate this process, they have given a problem to identify the customers segments, those are eligible for loan amount so that they can specifically target these customers. Here they have provided a data set.

Data

1. Variable Description
2. Loan ID - Unique Loan ID
3. Gender - Male/ Female
4. Married - Applicant married (Y/N)
5. Dependents - Number of dependents
6. Education - Applicant Education (Graduate/ Undergraduate)
7. Self Employed - Self-employed (Y/N)
8. Applicant Income - Applicant income
9. Co-applicant Income – Co-applicant income
10. Loan Amount - Loan amount in thousands
11. Loan Amount Term - Term of loan in months
12. Credit History - credit history meets guidelines
13. Property Area - Urban/ Semi Urban/ Rural
14. Loan Status - Loan approved (Y/N)

CODE: -

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: test_url = 'C:/Users/Saksham/Documents/Study Material/Internship/Finalprojects_DS-master/Loan_Prediction/test.csv'
train_url = 'C:/Users/Saksham/Documents/Study Material/Internship/Finalprojects_DS-master/Loan_Prediction/train.csv'
```

```
In [3]: train = pd.read_csv(train_url)
test = pd.read_csv(test_url)
```

```
In [4]: train.head()
```

```
Out[4]:
```


	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.0	1.0
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0

```
In [5]: plt.figure(6)

plt.subplot(121)
train_notnull = train.dropna()
sns.distplot(train_notnull['LoanAmount'])

plt.subplot(122)
train_notnull['LoanAmount'].plot.box(figsize = (16,5))
```


```
Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x20c8971f978>
```



```
In [6]: plt.figure(5)
plt.subplot(121)
sns.distplot(train['CoapplicantIncome'])

plt.subplot(122)
train['CoapplicantIncome'].plot.box(figsize = (16,5))
```

Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x20c8caa3eb8>




```
In [7]: plt.figure(2)
plt.subplot(231)
train['Dependents'].value_counts().plot.bar(figsize = (15,8), title = 'Dependents')

plt.subplot(232)
train['Property_Area'].value_counts().plot.bar(figsize = (15,8), title = 'Property_Area')

plt.subplot(233)
train['Education'].value_counts().plot.bar(figsize = (15,8), title = 'Education')
```

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x20c8c9fed30>



```
In [8]: train.corr()
```

Out[8]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
ApplicantIncome	1.000000	-0.116605	0.570909	-0.045306	-0.014715
CoapplicantIncome	-0.116605	1.000000	0.188619	-0.059878	-0.002056
LoanAmount	0.570909	0.188619	1.000000	0.039447	-0.008433
Loan_Amount_Term	-0.045306	-0.059878	0.039447	1.000000	0.001470
Credit_History	-0.014715	-0.002056	-0.008433	0.001470	1.000000

```
In [9]: train.columns
```

Out[9]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'], dtype='object')

```
In [10]: from sklearn import preprocessing
le = preprocessing.LabelEncoder()
```

```
In [11]: train['Gender'].unique()
```

Out[11]: array(['Male', 'Female', nan], dtype=object)

```
In [12]: train.isna().sum()
```

```
Out[12]: Loan_ID      0
Gender      13
Married     3
Dependents  15
Education   0
Self_Employed  32
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount   22
Loan_Amount_Term  14
Credit_History  50
Property_Area  0
Loan_Status  0
dtype: int64
```

```
In [13]: train = train[~train['Gender'].isna()]
```

```
In [14]: train = train[~train['Dependents'].isna()]
train = train[~train['Self_Employed'].isna()]
train = train[~train['LoanAmount'].isna()]
train = train[~train['Credit_History'].isna()]
train = train[~train['Loan_Amount_Term'].isna()]
```

```
In [15]: train.isna().sum()
```

```
Out[15]: Loan_ID      0
Gender      0
Married     0
Dependents   0
Education    0
Self_Employed  0
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount    0
Loan_Amount_Term  0
Credit_History  0
Property_Area  0
Loan_Status   0
dtype: int64
```

```
In [16]: from sklearn import preprocessing
le = preprocessing.LabelEncoder()
le.fit(train['Gender'])
x=le.transform(train['Gender'])
train['Gender'] = x
```

```
In [17]: le.fit(train['Married'])
x=le.transform(train['Married'])
train['Married'] = x
```

```
In [18]: le.fit(train['Dependents'])
x=le.transform(train['Dependents'])
train['Dependents'] = x
```

```
In [19]: le.fit(train['Education'])
x=le.transform(train['Education'])
train['Education'] = x
```

```
In [20]: le.fit(train['Self_Employed'])
x=le.transform(train['Self_Employed'])
train['Self_Employed'] = x
```

```
In [21]: le.fit(train['Property_Area'])
x=le.transform(train['Property_Area'])
train['Property_Area'] = x
```

```
In [22]: le.fit(train['Loan_Status'])
x=le.transform(train['Loan_Status'])
train['Loan_Status'] = x
```



```
In [21]: le.fit(train['Property_Area'])
x=le.transform(train['Property_Area'])
train['Property_Area'] = x
```

```
In [22]: le.fit(train['Loan_Status'])
x=le.transform(train['Loan_Status'])
train['Loan_Status'] = x
```

```
In [23]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 480 entries, 1 to 613
Data columns (total 13 columns):
Loan_ID      480 non-null object
Gender       480 non-null int32
Married      480 non-null int32
Dependents   480 non-null int32
Education    480 non-null int32
Self_Employed 480 non-null int32
ApplicantIncome 480 non-null int64
CoapplicantIncome 480 non-null float64
LoanAmount   480 non-null float64
Loan_Amount_Term 480 non-null float64
Credit_History 480 non-null float64
Property_Area 480 non-null int32
Loan_Status  480 non-null int32
dtypes: float64(4), int32(7), int64(1), object(1)
memory usage: 39.4+ KB
```

```
In [24]: train.head()
```

```
Out[24]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
1	LP001003	1	1	1	0	0	4583	1508.0	128.0	360.0	1.0
2	LP001005	1	1	0	0	1	3000	0.0	66.0	360.0	1.0
3	LP001006	1	1	0	1	0	2583	2358.0	120.0	360.0	1.0
4	LP001008	1	0	0	0	0	6000	0.0	141.0	360.0	1.0
5	LP001011	1	1	2	0	1	5417	4196.0	267.0	360.0	1.0

```
In [25]: cols = ['ApplicantIncome','CoapplicantIncome','LoanAmount','Loan_Amount_Term','Credit_History']
train[cols] = train[cols].applymap(np.int64)
```

```
In [26]: train.corr()
```

```
Out[26]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
Gender	1.000000	0.349424	0.217510	0.059245	-0.002761	0.032644	0.156170	0.098975	-0.088704	-0.000204
Married	0.349424	1.000000	0.386367	0.001652	0.015674	0.036717	0.102950	0.183442	-0.107504	-0.000204
Dependents	0.217510	0.386367	1.000000	0.028608	0.045754	0.131139	-0.000319	0.172780	-0.096361	-0.000204
Education	0.059245	0.001652	0.028608	1.000000	-0.005085	-0.131172	-0.074498	-0.172780	-0.102168	-0.000204
Self_Employed	-0.002761	0.015674	0.045754	-0.005085	1.000000	0.170785	-0.001508	0.120389	-0.034852	-0.000204
ApplicantIncome	0.032644	0.036717	0.131139	-0.131172	0.170785	1.000000	-0.112588	0.495310	-0.010838	-0.000204
CoapplicantIncome	0.156170	0.102950	-0.000319	-0.074498	-0.001508	-0.112588	1.000000	0.190740	-0.005773	-0.000204
LoanAmount	0.098975	0.183442	0.172780	-0.172780	0.120389	0.495310	0.190740	1.000000	0.050867	-0.000204
Loan_Amount_Term	-0.088704	-0.107504	-0.096361	-0.102168	-0.034852	-0.010838	-0.005773	0.050867	1.000000	-0.000204
Credit_History	0.022447	0.029095	-0.026651	-0.056656	-0.023568	-0.056152	-0.008692	-0.040773	0.032937	1.000000
Property_Area	-0.000204	0.038653	0.001191	-0.055005	-0.050797	-0.053160	0.006539	-0.109685	-0.058656	-0.000204
Loan_Status	0.064504	0.112321	0.035428	-0.068437	-0.034715	-0.043152	-0.049020	-0.071753	-0.007798	-0.000204

```
In [27]: x = train[['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History']]
```

```

In [28]: y = train[['Loan_Status']]

In [29]: from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size = 0.3)

In [30]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state = 0)

In [31]: model.fit(xtrain,ytrain)

C:\Users\Saksham\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in
0.22. Specify a solver to silence this warning.
  FutureWarning)
C:\Users\Saksham\Anaconda3\lib\site-packages\sklearn\utils\validation.py:761: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)

Out[31]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
  intercept_scaling=1, max_iter=100, multi_class='warn',
  n_jobs=None, penalty='l2', random_state=0, solver='warn',
  tol=0.0001, verbose=0, warm_start=False)

In [32]: model.score(xtest,ytest)

Out[32]: 0.8263888888888888

In [33]: model.score(xtrain,ytrain)

Out[33]: 0.7976190476190477

In [34]: from sklearn.ensemble import RandomForestClassifier
model_random = RandomForestClassifier(n_estimators = 60, max_depth = 1, random_state = 0, max_features = 7)

In [35]: model_random.fit(xtrain,ytrain)

C:\Users\Saksham\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: DataConversionWarning: A column-vector y was passed when a 1d array w
as expected. Please change the shape of y to (n_samples, ), for example using ravel().
  """Entry point for launching an IPython kernel.

Out[35]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
  max_depth=1, max_features=7, max_leaf_nodes=None,
  min_impurity_decrease=0.0, min_impurity_split=None,
  min_samples_leaf=1, min_samples_split=2,
  min_weight_fraction_leaf=0.0, n_estimators=60, n_jobs=None,
  oob_score=False, random_state=0, verbose=0, warm_start=False)

In [36]: model_random.score(xtrain,ytrain)

Out[36]: 0.7946428571428571

In [37]: model_random.score(xtest,ytest)

Out[37]: 0.8402777777777778

```

Accuracy: - 79.46%

Cross-Validation Score: - 84.02%